

Taiwan Semiconductor Manufacturing Co. (TSM)

Customer Economics & Revenue Concentration

A focused research note mapping TSMC's major customers, the products and manufacturing services they purchase, estimated annual spending, and the implications for the TSM investment thesis.

REPORT DATE August 10, 2026	TICKER NYSE: TSM / TWSE: 2330	RESEARCH VIEW Thesis-Supportive / Positive	SCOPE Customer & Revenue Exposure Note
--------------------------------	-------------------------------------	--	--

Bottom line: TSMC's customer concentration is high, but the underlying economic exposure is broader than the customer list suggests. NVIDIA and Apple represented approximately 36% of 2025 revenue, while the remaining top-eight customers represented another roughly \$51 billion. Importantly, TSMC participates across competing compute architectures: merchant GPUs, custom AI ASICs, CPUs, smartphones, networking and edge silicon. That makes customer concentration a real risk, but reduces the risk that a single architecture winner determines TSMC's long-term growth.

Executive Summary

TSMC reported US\$122.42 billion of 2025 revenue. Its ten largest customers generated approximately 78% of revenue, up from 76% in 2024. TSMC disclosed two customers above the 10% reporting threshold: “Customer A” at 19% of revenue and “Customer B” at 17%. Industry reporting and NVIDIA management commentary identify those customers as NVIDIA and Apple, respectively. This implies approximately \$23.3 billion of 2025 TSMC revenue from NVIDIA and \$20.8 billion from Apple.

The identities and spending of customers ranked below the top two are not disclosed by TSMC. The estimates in this report therefore use a constrained analyst model: the known \$122.42 billion revenue base, the 78% top-ten concentration, the disclosed NVIDIA and Apple amounts, confirmed customer relationships, product-node intensity, and the relative scale of each semiconductor program. The model is deliberately expressed as ranges rather than false-precision point estimates.


2025 TSMC REVENUE \$122.4B Company-reported, U.S. dollar basis	TOP 10 CUSTOMER CONCENTRATION 78% Approximately \$95.5B of 2025 revenue
NVIDIA + APPLE 36% Approximately \$44.1B combined	REMAINING TOP 8 ~\$51.4B Residual implied by TSMC disclosures


2025 Customer Revenue Map

Customer / Program	Primary Products / Services Purchased from TSMC	2025 Spend	TSMC Rev. Share	Evidence / Confidence
NVIDIA	Blackwell / Blackwell Ultra and Hopper-class AI GPUs; Grace CPU silicon; gaming / professional GPUs; custom N4/N4P process; CoWoS and related advanced packaging for HBM-based AI systems.	~\$23.3B	19%	Very High Amount disclosed as Customer A; identity strongly evidenced.
Apple	A-series iPhone/iPad SoCs, M-series Mac/iPad processors and other Apple-designed silicon; leading-edge 3nm/5nm/4nm production with migration toward N2; selected packaging/integration services.	~\$20.8B	17%	Very High Amount disclosed as Customer B; identity strongly evidenced.
Broadcom	Custom AI XPU / ASICs, data-center switching and networking silicon, high-speed connectivity; advanced nodes and 3D/2.5D packaging including TSMC-enabled multi-die AI designs.	\$8-10B	~6.5-8%	Medium Confirmed major TSMC customer; spend is analyst estimate.
AMD	EPYC CPUs, Ryzen CPUs/APUs, Radeon GPUs, Instinct AI accelerators and chiplets; N5/N4/N3-class wafers, advanced packaging, and N2 product migration including EPYC "Venice."	\$8-10B	~6.5-8%	Medium Confirmed major TSMC customer; spend is analyst estimate.
Qualcomm	Snapdragon mobile SoCs, Snapdragon X PC processors, automotive and connectivity silicon; substantial leading-edge 4nm/3nm fabrication, although not every Qualcomm product is necessarily sourced from TSMC.	\$7-9B	~5.5-7.5%	Medium Confirmed major TSMC customer; spend is analyst estimate.
MediaTek	Dimensity smartphone SoCs, connectivity, consumer/edge and emerging cloud/custom ASIC silicon; N3P and N2/N2P development and production.	\$6-8B	~5-6.5%	Medium Long-standing TSMC customer; spend is analyst estimate.
Intel	Externally manufactured supporting tiles/chiplets and selected client/graphics components; TSMC has historically supplied parts of Intel multi-tile products. Future spend could moderate as Intel ramps internal 18A.	\$5-7B	~4-6%	Medium-Low Relationship is established; 2025 scale is estimated and structurally variable.
Marvell	Custom AI accelerators, CPUs, networking/switching, optical DSP and data-infrastructure silicon; 3nm and 2nm platforms and CoWoS-enabled designs.	\$3-5B	~2.5-4%	Medium Advanced-node collaboration is public; spend is analyst estimate.
AWS / Annapurna or comparable hyperscaler program	Custom cloud CPUs and AI accelerators such as Graviton, Trainium and Inferentia-class silicon, either directly or through design/ASIC partners.	\$3-5B	~2.5-4%	Low-Medium Economic exposure is clear; direct invoiced-customer identity is not disclosed.
Other large customers	NXP, Realtek, Sony, automotive/industrial customers and other fabless or system-company programs; mix spans mature and advanced nodes.	\$3-5B	~2.5-4%	Low Residual bucket required to reconcile disclosed top-ten concentration.

Important: NVIDIA and Apple amounts are derived from company-disclosed customer shares. All other customer spending ranges are analyst estimates, not TSMC-reported figures. Google, Meta and Microsoft are significant economic end-demand drivers but may purchase TSMC capacity indirectly through NVIDIA, Broadcom, Marvell, AMD or other design partners; counting them separately can double-count the same TSMC revenue.

Illustrative 2025 Customer Spend





Customer Profiles & Strategic Importance

1. NVIDIA - AI Infrastructure Anchor AI/HPC 4NP CoWoS WHAT TSMC PROVIDES NVIDIA's Blackwell architecture uses a custom-built TSMC 4NP process. TSMC also provides the advanced packaging ecosystem necessary to combine large GPU dies with HBM at scale. That means TSMC monetizes both front-end wafer fabrication and a growing amount of back-end system integration.	~\$23.3B INVESTMENT SIGNIFICANCE NVIDIA became TSMC's largest customer in 2025, illustrating how AI has changed TSMC's revenue mix. The relationship is powerful but creates concentration risk: a major AI capex correction, inventory digestion, or foundry diversification by NVIDIA would be material.
2. Apple - Leading-Edge Volume Anchor Smartphone PC N3 / N2 WHAT TSMC PROVIDES TSMC manufactures Apple-designed silicon across iPhone, iPad and Mac product families. Apple has been a critical early adopter of advanced nodes and remains the first and largest customer of TSMC's Arizona fab; Apple said it expected to buy well over 100 million advanced chips from that facility in 2026.	~\$20.8B INVESTMENT SIGNIFICANCE Apple provides high-volume, predictable leading-edge utilization and has historically helped de-risk new-node ramps. Its share of TSMC revenue fell as AI outgrew smartphones, but absolute spending remains enormous. The key risk is diversification toward alternative foundries or slower device unit growth.
3. Broadcom - Custom AI and Networking Growth Engine Custom XPU Networking 3.5D WHAT TSMC PROVIDES Broadcom uses TSMC advanced nodes and packaging for custom AI accelerators and large networking ASICs. Broadcom has publicly described AI XPU designs using TSMC process technology and advanced face-to-face / multi-die integration.	\$8-10B est. INVESTMENT SIGNIFICANCE Broadcom is central to the "customers behind the customers" thesis. Hyperscalers can shift AI spend away from merchant GPUs toward custom accelerators without necessarily shifting foundry economics away from TSMC. Broadcom therefore partially hedges TSMC against NVIDIA market-share loss.
4. AMD - Broad Compute Exposure CPU GPU AI N2 WHAT TSMC PROVIDES AMD relies on TSMC for EPYC server CPUs, Ryzen client processors, Radeon GPUs and Instinct AI accelerators. AMD's next-generation EPYC "Venice" achieved the first HPC product silicon milestone on TSMC N2, underscoring the depth of the roadmap relationship.	\$8-10B est. INVESTMENT SIGNIFICANCE AMD gives TSMC exposure to both conventional CPU share gains and AI accelerator growth. AMD's chiplet architecture also increases the strategic value of advanced packaging and heterogeneous integration.

5. Qualcomm - Mobile, PC and Automotive Scale

\$7-9B est.

Mobile PC Auto 3nm

WHAT TSMC PROVIDES

Qualcomm is a long-standing TSMC customer. Its premium Snapdragon platforms increasingly use leading-edge process nodes, while the broader portfolio spans mobile SoCs, Windows-on-Arm PC processors, automotive compute and connectivity.

INVESTMENT SIGNIFICANCE

Qualcomm supplies scale outside AI data centers, helping diversify TSMC's end markets. The relationship is somewhat less captive than NVIDIA or AMD because Qualcomm has historically used multiple foundries across product generations.

6. MediaTek - High-Volume Edge Silicon Moving to N2

\$6-8B est.

Smartphone Edge AI N3P N2P

WHAT TSMC PROVIDES

MediaTek's flagship Dimensity 9500 uses TSMC N3P, and MediaTek has publicly announced development on TSMC 2nm/N2P. Its product set spans smartphones, Wi-Fi/connectivity, consumer electronics, PCs and emerging custom AI / data-center silicon.

INVESTMENT SIGNIFICANCE

MediaTek is both a volume anchor and an important N2 migration customer. Its expansion into custom ASIC and cloud silicon could increase TSMC content per customer beyond traditional smartphones.

7. Intel - Material but Strategically Variable Outsourcing

\$5-7B est.

Tiles / Chiplets Client Foundry Diversification

WHAT TSMC PROVIDES

Intel has used TSMC for supporting tiles and externally sourced components in multi-tile products. Outsourcing grew as Intel sought access to competitive external process technology and flexible capacity.

INVESTMENT SIGNIFICANCE

Unlike fabless customers, Intel has a strategic incentive to bring more leading-edge production back in-house as 18A scales. Intel should therefore be treated as a less durable source of TSMC wallet share even if near-term spending remains meaningful.

8. Marvell - Custom Silicon and Interconnect Lever

\$3-5B est.

Custom ASIC Networking Optical 2nm

WHAT TSMC PROVIDES

Marvell has publicly demonstrated 2nm silicon on TSMC for next-generation custom AI accelerators, CPUs and networking, and has extended its TSMC collaboration across 3nm, 2nm and CoWoS-enabled data infrastructure products.

INVESTMENT SIGNIFICANCE

Marvell represents the long-tail upside from custom cloud silicon. As hyperscalers diversify accelerators and networking, TSMC can monetize more designs even when the end customer is not directly visible in TSMC's disclosures.

Customers Behind the Customers

A critical analytical distinction is the difference between **TSMC's invoiced customer** and the **ultimate economic buyer of compute**. Google, Microsoft, Meta, Amazon and other cloud operators can influence TSMC demand through multiple routes:

Merchant GPU route

Cloud provider buys NVIDIA or AMD accelerators -> NVIDIA/AMD buys wafers and advanced packaging from TSMC.

Custom accelerator route

Cloud provider develops a custom accelerator with Broadcom, Marvell or another ASIC partner -> design partner and/or hyperscaler books TSMC capacity.

Custom CPU route	Cloud provider develops Arm/RISC-V CPU silicon -> TSMC captures the foundry economics regardless of whether the CPU competes with x86.
Networking route	Higher AI cluster scale drives switch, NIC, optical DSP and interconnect silicon -> additional TSMC demand through Broadcom, Marvell and other networking customers.

Analyst interpretation: TSMC is not merely exposed to NVIDIA's share of AI compute. It is exposed to the overall silicon intensity of AI infrastructure. If a hyperscaler shifts dollars from NVIDIA GPUs to a Broadcom-designed custom XPU, part of the competitive loss for NVIDIA can remain a revenue opportunity for TSMC.

2026 Evidence: The Mix Is Moving Further Toward HPC

TSMC's 2Q26 results reinforce the customer thesis. HPC accounted for **66% of quarterly revenue**, up 20% sequentially, while smartphones accounted for 22%. Advanced technologies at 7nm and below represented **77% of wafer revenue**, and 2nm contributed 3% in its early ramp. Management raised full-year 2026 revenue growth expectations to slightly above 40% in U.S. dollar terms and increased planned capital spending to \$60-\$64 billion.

On August 10, TSMC reported July 2026 revenue of NT\$467.58 billion, up 44.7% year over year; January-through-July revenue rose 37.0%. These figures are consistent with sustained leading-edge and AI-driven customer demand rather than a one-quarter spike.

Investment Thesis Implications

What is working

- **Wallet-share expansion:** TSMC increasingly sells both leading-edge wafers and advanced packaging, raising revenue captured per AI accelerator program.
- **Architecture-neutral exposure:** NVIDIA, AMD, Broadcom, Marvell and hyperscaler custom silicon compete with one another, but all can still expand TSMC's addressable revenue.
- **Node migration:** N2 is already contributing revenue, with major customers such as AMD and MediaTek publicly advancing products on the node.
- **Capacity visibility:** TSMC says it works with customers and customers' customers on multi-year roadmaps, supporting unusually long demand visibility for advanced-node investment.
- **Customer depth:** The top-ten concentration reflects technological dependence as much as financial concentration. Switching leading-edge foundries requires years of design, qualification and ramp work.

What could weaken the thesis

- **Concentration rises further:** NVIDIA alone at 19% is already material; a move toward 25%+ would make TSMC more dependent on one AI capex cycle and one customer's execution.
- **Foundry diversification:** Successful Samsung or Intel foundry ramps that win meaningful leading-edge share from Apple, NVIDIA, Qualcomm or other customers would impair TSMC's pricing power and utilization.
- **AI capex synchronization:** TSMC is diversified across AI architectures, but not against a broad decline in hyperscaler AI infrastructure spending. A simultaneous slowdown would hit NVIDIA, custom ASICs, networking and packaging together.
- **Intel insourcing:** Intel's 18A ramp can reduce external foundry spending even if Intel's end-product volumes recover.
- **Geopolitical/export controls:** Restrictions on advanced chips can change customer product mixes and create compliance risk around indirect end users.

Key Monitoring Framework

Checkpoint	Why It Matters	Constructive Signal	Warning Signal
NVIDIA share of TSMC revenue	Measures concentration and AI dependence.	Growth with top-ten diversification.	NVIDIA rises materially above ~20% while other customers stagnate.
Broadcom / Marvell custom silicon growth	Tests the “TSMC wins either way” custom-ASIC thesis.	More 3nm/2nm XPU programs and advanced packaging demand.	Custom silicon migrates to competing foundries.
Apple N2 / future-node allocation	Apple remains the largest smartphone/ consumer anchor.	TSMC retains dominant Apple leading-edge share.	Meaningful Apple allocation shifts to Intel/Samsung.
CoWoS / advanced packaging capacity	Advanced packaging is an increasingly important AI bottleneck and revenue layer.	Capacity expands with stable pricing and utilization.	Oversupply, major price compression or competitive share loss.
N2/A16 customer adoption	Determines next wave of wafer ASP and gross-profit growth.	Broad customer tape-outs across AI, CPUs and mobile.	Delayed ramps, yield issues or weak customer conversion.
HPC revenue share and growth	Best reported proxy for AI/compute customer demand.	High growth with broad customer participation.	Sequential HPC contraction combined with lower utilization.

Conclusion

The customer analysis is **constructive for the TSM investment thesis**. The obvious concern is concentration: two customers generated about 36% of 2025 revenue and the top ten generated 78%. But the composition of those customers is unusually attractive. TSMC's largest buyers are the companies driving the highest-value semiconductor workloads in AI, smartphones, CPUs, networking and custom cloud silicon.

The stronger insight is that TSMC's economic diversification is better than its accounting concentration suggests. NVIDIA and Broadcom compete in AI accelerators; AMD competes with NVIDIA and Intel; Apple competes in PCs and smartphones; Qualcomm and MediaTek compete in mobile; hyperscalers increasingly design their own chips. Yet each competitive branch can still lead back to TSMC's advanced-node and packaging ecosystem.

Research view: Customer concentration should remain a monitored risk, but it is not currently a thesis-breaker. The more important thesis-break condition would be evidence that major leading-edge customers can shift meaningful volumes to another foundry without sacrificing yield, performance, time-to-market or packaging integration. Until that occurs, TSMC's customer roster remains one of the clearest expressions of its technological moat.

Methodology & Limitations

- TSMC publicly discloses only customers representing at least 10% of annual revenue. For 2025, those were anonymized as Customer A (19%) and Customer B (17%).
- NVIDIA and Apple identities are inferred from public reporting and management statements; the underlying revenue amounts are TSMC-disclosed.
- All spending estimates below the top two are model ranges, not company guidance or audited customer-level revenue.
- The model reconciles to the top-ten residual: approximately \$95.5B of top-ten revenue less ~\$44.1B from NVIDIA and Apple leaves roughly \$51.4B attributable to the other eight largest customers.
- Ultimate cloud end-demand is not the same as TSMC billing identity. Google, Microsoft, Meta and Amazon may drive TSMC revenue indirectly through NVIDIA, AMD, Broadcom, Marvell or other ASIC partners.
- This report is a customer-economics research note, not a full valuation update or price-target report.

Selected Sources

1. TSMC 2025 Form 20-F, U.S. SEC - top-ten customer concentration of 78%. SEC filing.
2. TSMC 2025 consolidated financial report - Customer A NT\$726.974B / 19%; Customer B NT\$645.179B / 17%; 2025 platform and node revenue. SEC filing.
3. TSMC 2025 Annual Report - US\$122.42B consolidated revenue. TSMC Investor Relations.
4. Taiwan News/CNA, Feb. 27, 2026 - identifies NVIDIA as Customer A and Apple as Customer B, citing NVIDIA CEO Jensen Huang and TSMC's financial report. Taiwan News.
5. TSMC 2Q26 earnings call and presentation - HPC 66% of revenue; 7nm-and-below 77% of wafer revenue; N2 3%; 2026 revenue growth slightly above 40%; capex \$60-\$64B. TSMC Investor Relations.
6. TSMC July 2026 Revenue Report - NT\$467.58B, +44.7% YoY; Jan-Jul +37.0%. TSMC Press Center.
7. TSMC U.S. investment announcement naming Apple, NVIDIA, AMD, Broadcom and Qualcomm among leading U.S. customers. TSMC Press Center.
8. NVIDIA technical documentation - Blackwell GPUs manufactured using custom TSMC 4NP process. NVIDIA.
9. TSMC CoWoS - advanced 2.5D packaging for HPC/AI integrating SoCs and HBM. TSMC 3DFabric.
10. Apple - TSMC Arizona relationship; Apple is the facility's first/largest customer and expects to purchase well over 100M advanced chips from Arizona in 2026. Apple Newsroom.
11. AMD - next-generation EPYC "Venice" is AMD's first HPC product on TSMC N2. AMD Newsroom.
12. Broadcom - AI XPU multi-die architecture leveraging TSMC process nodes and advanced integration. Broadcom Investor Relations.
13. MediaTek - Dimensity 9500 on TSMC N3P and 2nm/N2P product development. MediaTek; MediaTek 2nm release.
14. Marvell - 2nm silicon on TSMC for custom AI accelerators, CPUs and networking. Marvell.
15. Intel - prior disclosure that supporting tiles for multi-tile products were manufactured at TSMC, illustrating the outsourcing relationship. Intel Newsroom archive.

Prepared as an independent analytical research note for informational purposes. Estimates are explicitly identified and should not be treated as customer disclosures, audited data, or investment advice. Dollar conversions and customer-share calculations are approximate due to exchange-rate and rounding differences.